

# TTIC 31230, Fundamentals of Deep Learning

David McAllester, Winter 2019

## Connectionist Temporal Classification (CTC)

and Deep Graphical Models

## The Fundamental Equation: Conditional vs. Unconditional

$$\Phi^* = \operatorname{argmin}_{\Phi} E_{(x,y) \sim \text{Pop}} - \ln P(y|x)$$

$$\Phi^* = \operatorname{argmin}_{\Phi} E_{y \sim \text{Pop}} - \ln P(y)$$

This is a non-distinction: the analysis of the conditional case is exactly the same as that of the unconditional case.

# The Fundamental Equation:

## Distributions on Exponentially Large Sets

The structured case:  $y \in \mathcal{Y}$  where  $\mathcal{Y}$  is discrete but iteration over  $\hat{y} \in \mathcal{Y}$  is infeasible.

Language modeling (unconditional) and machine translation (conditional) are distributions on exponentially large (even infinite) sets.

## Friendly and Unfriendly Distributions

A model  $P_{\Phi}(y)$  will be called **friendly** if we can efficiently sample from it and, for any given  $y$ , can efficiently compute  $P_{\Phi}(y)$ .


**Autoregressive** language models (unconditional) and autoregressive machine translation models (conditional) are **friendly**.

Distributions which are not friendly in this sense will be called **unfriendly**.

## The Importance of Being Friendly

If  $P_{\Phi}(y|x)$  can be computed (a friendly model) we can do SGD on cross-entropy loss  $-\ln P_{\Phi}(y|x)$  by back-propagating through the computation of  $P_{\Phi}(y|x)$ .

# Graphical Models



Koller and Friedman, MIT Press, 2009, 1270 pages

# Semantic Segmentation




SLIC superpixels, Achanta et al.

We want to assign each superpixel one of  $k$  semantic classes.

For example “person”, “car”, “building”, “sky” or “other”.

# General Markov Random Fields (MRFs)



$\hat{y}$  assigns a class  $\hat{y}[i]$  to each node (superpixel)  $i$ .

$$s(\hat{y}) = \sum_{i \in \text{Nodes}} s_i[\hat{y}[i]] + \sum_{e \in \text{Edges}} s_e[\hat{y}[e.i], \hat{y}[e.j]]$$

Node Potentials

Edge Potentials



## An Example

Consider an image with three superpixels  $A$ ,  $B$  and  $C$  where each superpixel is to be labeled as either “foreground” or background.

Suppose the unary potentials are all zero.

$$s_A(\text{Foreground}) = s_A(\text{Background}) = 0$$

$$s_B(\text{Foreground}) = s_B(\text{Background}) = 0$$

$$s_C(\text{Foreground}) = s_C(\text{Background}) = 0$$

## The Binary Potentials

Let  $F_A$  be the proposition that  $A$  is foreground and similarly for  $F_B$  and  $F_C$ .

We can express  $P_A \Rightarrow P_B$  with

$$s_{A,B}(\text{Foreground}, \text{Background}) = -1$$

$$s_{A,B}(\text{Foreground}, \text{Foreground}) = 1$$

$$s_{A,B}(\text{Background}, \text{Background}) = 1$$

$$s_{A,B}(\text{Background}, \text{Foreground}) = 1$$

The binary potentials are then given by  $F_A \Rightarrow F_B$ ,  $F_B \Rightarrow F_C$ ,  $F_C \Rightarrow F_A$ .

## The Full Configuration Potential

For any configuration  $\hat{y}$  we have that  $s(\hat{y})$  is the sum of the unary and binary potentials.

If none are foreground we have  $s(\hat{y}) = 3$


If one is foreground we have  $s(\hat{y}) = -1 + 1 + 1 = 1$

If two are foreground we also have  $s(\hat{y}) = -1 + 1 + 1 = 1$

If all are foreground we have  $s(\hat{y}) = 3$ .

$$Z = 6 * 1 + 2 * 3 = 12 \quad P_A(\text{Foreground}) = \frac{3 * 1 + 3}{12} = \frac{1}{2}$$

# Exponential Softmax




$\hat{y}$  assigns a class  $\hat{y}[i]$  to each node (superpixel)  $i$ .

$$P_{\Phi}(\hat{y}|x) = \underset{\hat{y}}{\text{expsoftmax}} \ s_{\Phi}(\hat{y}|x)$$

$$s_{\Phi}(\hat{y}|x) = \sum_{i \in \mathcal{I}} s_i[\hat{y}[i]] + \sum_{e \in \mathcal{E}} s_e[\hat{y}[e.i], \hat{y}[e.j]]$$

## Exponential Softmax is Typically Unfriendly



$\hat{y}$  assigns a class  $\hat{y}[i]$  to each node (superpixel)  $i$ .

$$P_{\Phi}(\hat{y}|x) = \underset{\hat{y}}{\text{expsoftmax}} \ s_{\Phi}(\hat{y}|x)$$

Computing  $Z$  in general is #P hard.

But special cases can be friendly and approximations can be made in unfriendly cases.

## Latent Variables

We are often interested in models of the form

$$P_{\Phi}(y) = \sum_z P_{\Phi}(z) P_{\Phi}(y|z).$$

Probabilistic grammar models have this form where  $y$  is a sentence and  $z$  is a parse tree and  $P(y|z)$  is deterministic.

## Exponential Softmax as Intermediate Computation

$$P_{\Phi}(y) = \sum_z P_{\Phi}(z) P_{\Phi}(y|z).$$

input  $x$

$\vdots$

$z = \text{expsoftmax}_{\mathbf{z}} \dots$

input  $z$

$\vdots$

$y = \text{expsoftmax}_{\mathbf{y}} \dots$

## A Composition of Friendlies is Typically Unfriendly

$$P_{\Phi}(y) = \sum_z P_{\Phi}(z)P_{\Phi}(y|z).$$

It is often the case that  $P_{\Phi}(z)$  is friendly, and  $P_{\Phi}(y|z)$  is friendly, but  $P_{\Phi}(y)$  is not friendly (the sum over  $z$  is intractible).

For example  $z$  might be uniformly distributed over assignments of truth values to Boolean variables (which is friendly) and  $y$  might be the value of a fixed Boolean formula  $\Phi$  (which is friendly given  $z$ ). In this case determining if  $P_{\Phi}(y) > 0$  is the SAT problem which is NP hard.



# Connectionist Temporal Classification (CTC)

## A Successful Deep Latent Variable Model

A speech signal

$$x = x_1, \dots, x_T$$

is labeled with a phone sequence

$$y = y_1, \dots, y_N$$

with  $N \ll T$  and with  $y_n \in \mathcal{P}$  for a set of phonemes  $\mathcal{P}$ .

The length  $N$  of  $y$  is not determined by  $x$  and the alignment between  $x$  and  $y$  is not given.

## CTC: A Friendly Compositions of Friendlies

$$P_{\Phi}(y|x) = \sum_z P_{\Phi}(z|x)P_{\Phi}(y|z).$$

Input Signal:  $x = x_1, \dots, x_T$

Latent Label:  $z = z_1, \dots, z_T, \quad z_t \in \mathcal{P} \cup \{\perp\}$

Output:  $y(z) = y_1, \dots, y_N$

$y(z)$  is the result of removing all the occurrences of  $\perp$  from  $z$ :

$$z \Rightarrow y$$

$$\perp, a_1, \perp, \perp, \perp, a_2, \perp, \perp, a_3, \perp \Rightarrow a_1, a_2, a_3$$

## The CTC Model

For  $z \in \mathcal{P} \cup \{\perp\}$  we have an embedding  $e(z)$ . The embedding is a parameter of the model.

$$h_1, \dots, h_T = \text{RNN}_\Phi(x_1, \dots, x_T)$$

$$P_\Phi(z_t | x_1, \dots, x_T) = \underset{z}{\text{softmax}} \ e(z)^\top h_t$$

$z_1, \dots, z_T$  are **all independent** given  $x$  (very friendly).

$P_\Phi(y|z)$  is **deterministic** (very friendly).

But it is not obvious whether  $P_\Phi(y|x)$  is friendly.

# Dynamic Programming

$$x = x_1, \dots, x_T$$

$$z = z_1, \dots, z_T, \quad z_t \in \mathcal{P} \cup \{\perp\}$$

$$y = y_1, \dots, y_N, \quad y_n \in \mathcal{P}, \quad N \ll T$$

$$y(z) = (z_1, \dots, z_T) - \perp$$

$$\vec{y}_t = (z_1, \dots, z_t) - \perp$$

$$\textcolor{red}{F}[\textcolor{red}{n}, \textcolor{red}{t}] = P(\vec{y}_{\textcolor{red}{t}} = y_1, \dots, y_{\textcolor{red}{n}})$$

$$P(y) = F[N, T]$$

# Dynamic Programming

$$\vec{y}_t = (z_1, \dots, z_t) - \perp$$
$$F[n, t] = P(\vec{y}_t = y_1, \dots, y_n)$$

$$F[0, 0] = 1$$

$$F[n, 0] = 0 \quad \text{for } n > 0$$

$$F[n + 1, t + 1] = P(z_{t+1} = \perp)F[n + 1, t] + P(z_{t+1} = y_{n+1})F[n, t]$$

# Semantic Segmentation




SLIC superpixels, Achanta et al.

We want to assign each superpixel one of  $k$  semantic classes.

For example “person”, “car”, “building”, “sky” or “other”.

## Unfriendly Exponential Softmax



$\hat{y}$  assigns a class  $\hat{y}[i]$  to each node (superpixel)  $i$ .

$$P_{\Phi}(\hat{y}|x) = \underset{\hat{y}}{\text{expsoftmax}} \ s_{\Phi}(\hat{y}|x)$$

$$s_{\Phi}(\hat{y}|x) = \sum_{i \in \mathcal{I}} s_i[\hat{y}[i]] + \sum_{e \in \mathcal{E}} s_e[\hat{y}[e.i], \hat{y}[e.j]]$$

# Back-Propagation Through Unfriendly Softmax

input  $x$

$\vdots$

$$s_i[c] = \dots$$

$$s_e[c, c'] = \dots$$


$$\mathcal{L} = -\ln P(y \mid s_{\mathcal{I}}[\mathcal{C}], s_{\mathcal{E}}[\mathcal{C}, \mathcal{C}])$$

We need to compute  $s_i.\text{grad}[c]$  and  $s_e.\text{grad}[c, c']$ .



# Hyper-Graphs: More General and More Concise

A hyper-edge is a subset of nodes.



$$s(\hat{y}) = \sum_{i \in \text{Nodes}} s_i[\hat{y}[i]] + \sum_{e \in \text{Edges}} s_e[\hat{y}[e.i], \hat{y}[e.j]]$$

$$s(\hat{y}) = \sum_{e \in \text{HyperEdges}} s_e[\hat{y}[e]]$$

## Backpropagation

The input is the image  $x$  and the parameter package  $\Phi$

$$\begin{aligned} & \vdots \\ s_e[\hat{y}] &= \dots \\ \mathcal{L} &= -\ln P(y \mid s_{\mathcal{E}}[\mathcal{Y}]) \end{aligned}$$

We abbreviate  $P(\hat{y} \mid s_{\mathcal{E}}[\mathcal{Y}])$  as  $P_s(\hat{y})$  — the distribution on  $\hat{y}$  defined by the tensor  $s$ .

We need to compute  $\nabla_s -\ln P_s(y)$ , or equivalently,  $s_e.\text{grad}[\hat{y}[e]]$ .

## Back-Propagation Through An Exponential Softmax

We will abbreviate  $s_e[\hat{y}[e]]$  as  $s_e[\tilde{y}]$ .

$\tilde{y}$  has a small number of possible values.

We will similarly write  $s_e.\text{grad}[\tilde{y}]$ .

We need to compute the tensor values  $s_e.\text{grad}[\tilde{y}]$

## Back-Propagation Through An Exponential Softmax

$$\begin{aligned}\text{loss}(s, y) &= -\ln \left( \frac{1}{Z(s)} e^{s(y)} \right) \\ &= \ln Z(s) - s(y)\end{aligned}$$

$$s_e.\text{grad}[\tilde{y}] = \left( \frac{1}{Z} \sum_{\hat{y}} e^{s(\hat{y})} (\partial s(\hat{y}) / \partial s_e[\tilde{y}]) \right) - (\partial s(y) / \partial s_e[\tilde{y}])$$

## Back-Propagation Through An Exponential Softmax

$$\begin{aligned} s_e.\text{grad}[\tilde{y}] &= \left( \frac{1}{Z} \sum_{\hat{y}} e^{s(\hat{y})} (\partial s(\hat{y}) / \partial s_e[\tilde{y}]) \right) - (\partial s(y) / \partial s_e[\tilde{y}]) \\ &= \left( \sum_{\hat{y}} P_s(\hat{y}) (\partial s(\hat{y}) / \partial s_e[\tilde{y}]) \right) - (\partial s(y) / \partial s_e[\tilde{y}]) \\ &= E_{\hat{y} \sim P_s} \mathbb{1}[\hat{y}[e] = \tilde{y}] - \mathbb{1}[y[e] = \tilde{y}] \\ &= P_{\hat{y} \sim P_s}(\hat{y}[e] = \tilde{y}) - \mathbb{1}[y[e] = \tilde{y}] \end{aligned}$$


## Hyperedge Marginals

$$s.\text{grad}[e, \tilde{y}] = P_{\hat{y} \sim P_s}(\hat{y}[e] = \tilde{y}) - \mathbb{1}[y[e] = \tilde{y}]$$

We write  $P_e(\tilde{y})$  for the hyperedge marginal  $P_{\hat{y} \sim P_s}(\hat{y}(e) = \tilde{y})$ .

To back-propagate log loss on a labeling of an unfriendly MRF it suffices to compute (or perhaps approximate) the current model's hyperedge marginals  $P_e(\tilde{y})$ .

## Tree-Structured Models are Friendly




Tree structure models can always be locally renormalized to form “autoregressive” models that predict one node at a time.

Also, the hyperedge marginals can be computed exactly.

$$s.\text{grad}[e, \tilde{y}] = P_e(\tilde{y}) - \mathbb{1}[y[e] = \tilde{y}]$$

## Belief Propagation



Belief Propagation is a message passing procedure (actually dynamic programming).


For each edge  $\{i, j\}$  and possible value  $\tilde{y}$  for node  $i$  we define  $Z_{j \rightarrow i}[\tilde{y}]$  to be the partition function for the subtree attached to  $i$  through  $j$  and with  $\hat{y}[i]$  restricted to  $\tilde{y}$ .

The function  $Z_{j \rightarrow i}$  on the possible values of node  $i$  is called the **message** from  $j$  to  $i$ .

The reverse direction message  $Z_{i \rightarrow j}$  is defined similarly.




## Computing the Messages



$$Z_{j \rightarrow i}[\tilde{y}] = \sum_{\tilde{y}'} e^{s_j[\tilde{y}'] + s_{\{j,i\}}[\{\tilde{y}', \tilde{y}\}]} \left( \prod_{k \in N(j), k \neq i} Z_{k \rightarrow j}[\tilde{y}'] \right)$$

## Computing Node Marginals from Messages



$$\begin{aligned} Z_i(\tilde{y}) &\doteq \sum_{\hat{y}: \hat{y}[i]=\tilde{y}} e^{s(\hat{y})} \\ &= e^{s_i[\tilde{y}]} \left( \prod_{j \in N(i)} Z_{j \rightarrow i}[\tilde{y}] \right) \\ \textcolor{red}{P}_i(\tilde{y}) &= Z_i(\tilde{y})/Z, \quad Z = \sum_{\tilde{y}} Z_i(\tilde{y}) \end{aligned}$$

## Computing Edge Marginals from Messages

$$\begin{aligned} Z_{\{i,j\}}(\tilde{y}) &\doteq \sum_{\hat{y}: \hat{y}[\{i,j\}]=\tilde{y}} e^{s(\hat{y})} \\ &= e^{s[i,\tilde{y}[i]]+s[j,\tilde{y}[j]]+s[\{i,j\},\tilde{y}]} \\ &\quad \prod_{k \in N(i), k \neq j} Z_{k \rightarrow i}[\tilde{y}[i]] \\ &\quad \prod_{k \in N(j), k \neq i} Z_{k \rightarrow j}[\tilde{y}[j]] \end{aligned}$$

$$P_{\{i,j\}}(\tilde{y}) = Z_{\{i,j\}}(\tilde{y})/Z$$

## Loopy BP

Message passing is also called belief propagation (BP).

In a graph with cycles it is common to do **Loopy BP**.

This is done by initializing all message  $Z_{i \rightarrow j}[\tilde{y}] = 1$  and then repeating (until convergence) the updates

$$P_{j \rightarrow i}[\tilde{y}] = \frac{1}{Z} Z_{j \rightarrow i}[\tilde{y}] \quad Z = \sum_{\tilde{y}} Z_{j \rightarrow i}[\tilde{y}]$$

$$Z_{j \rightarrow i}[\tilde{y}] = \sum_{\tilde{y}'} e^{s[j, \tilde{y}'] + s[\{j, i\}, \{\tilde{y}', \tilde{y}\}]} \left( \prod_{k \in N(j), k \neq i} P_{k \rightarrow j}[\tilde{y}'] \right)$$

# Other Methods of Approximating Hyperedge Marginals

MCMC Sampling

Contrastive Divergence

Pseudo-Likelihood

# Sampling

The quantities  $P_e(\tilde{e})$  are **hyperedge marginals**.

We can estimate the hyperedge marginals by sampling  $\hat{y}$  from  $P_s(\hat{y})$ .

# Monte Carlo Markov Chain (MCMC) Sampling

## Metropolis Algorithm

Pick an initial graph label  $\hat{y}$  and then repeat:

1. Pick a “neighbor”  $\hat{y}'$  of  $\hat{y}$  uniformly at random. The neighbor relation must be symmetric. Perhaps Hamming distance one.
2. If  $s(\hat{y}') > s(\hat{y})$  update  $\hat{y} = \hat{y}'$
3. If  $s(\hat{y}') \leq s(\hat{y})$  then update  $\hat{y} = \hat{y}'$  with probability  $e^{-(s(\hat{y}) - s(\hat{y}'))}$

## Markov Processes and Stationary Distributions

A Markov process is a process defined by a fixed state transition probability  $P(\hat{y}'|\hat{y}) = M_{\hat{y}',\hat{y}}$ .

Let  $P^t$  the probability distribution for time  $t$ .

$$P^{t+1} = MP^t$$

If every state can be reached from every state (ergodic process) then  $P^t$  converges to a unique **stationary distribution**  $P^\infty$

$$P^\infty = MP^\infty$$



## Metropolis Theorem

To verify that the Metropolis process has the correct stationary distribution we simply verify that  $MP = P$  where  $P$  is the desired distribution.

This can be done by checking that under the desired distribution the flow from  $\hat{y}$  to  $\hat{y}'$  equals the flow from  $\hat{y}'$  to  $\hat{y}$  (**detailed balance**).

## Metropolis Theorem

For  $s(\hat{y}) \geq s(\hat{y}')$

$$\text{flow}(\hat{y}' \rightarrow \hat{y}) = \frac{1}{Z} e^{s(\hat{y}')} \frac{1}{N}$$

$$\text{flow}(\hat{y} \rightarrow \hat{y}') = \frac{1}{Z} e^{s(\hat{y})} \frac{1}{N} e^{-\Delta f} = \frac{1}{Z} e^{s(\hat{y}')} \frac{1}{N}$$

But detailed balance is not required in general (see Hamiltonian MCMC).

## Gibbs Sampling

The Metropolis algorithm wastes time by rejecting proposed moves.

Gibbs sampling avoids this move rejection.

In Gibbs sampling we select a node  $i$  at random and change that node by drawing a new node value conditioned on the current values of the other nodes.

## Gibbs Sampling

$$P_s(i = \tilde{y} \mid \hat{y}) \doteq P_s(\hat{y}[i] = \tilde{y} \mid \hat{y}[1], \dots, \hat{y}[i-1], \hat{y}[i+1], \dots, \hat{y}[I])$$

Markov Blanket Property:

$$P_s(i = \tilde{y} \mid \hat{y}) = P_s(i = \tilde{y} \mid \hat{y}[N(i)])$$

Gibbs Sampling, Repeat:

- Select  $i$  at random
- draw  $\tilde{y}$  from  $P_s(i = \tilde{y} \mid \hat{y})$
- $\hat{y}[i] = \tilde{y}$

## Gibbs Sampling

Let  $\hat{y}[i = \tilde{y}]$  be the assignment  $\hat{y}'$  equal to  $\hat{y}$  except  $\hat{y}'[i] = \tilde{y}$ .

$$\begin{aligned} P_s(i = \tilde{y} \mid \hat{y}) &= \frac{P_s(\hat{y}[i] = \tilde{y})}{\sum_{\tilde{y}} P_s(\hat{y}[i] = \tilde{y})} \\ &= \frac{e^{s(\hat{y}[i=\tilde{y}])}}{\sum_{\tilde{y}} e^{s(\hat{y}[i=\tilde{y}])}} \end{aligned}$$

## Gibbs Sampling Theorem

$P_s(\hat{y})$  is a stationary distribution of Gibbs Sampling.

- Select  $i$  at random
- draw  $\tilde{y}$  from  $P_s(i = \tilde{y} \mid \hat{y})$
- $\hat{y}[i] = \tilde{y}$

The distribution before the update equals the distribution after the update.

## Pseudolikelihood

In Pseudolikelihood we replace the objective  $-\log P_s(\hat{y})$  with the objective  $-\log \tilde{Q}_s(\hat{y})$  where

$$\tilde{Q}_s(\hat{y}) \doteq \prod_i P_s(i = \hat{y}[i] \mid \hat{y})$$

$$\text{loss}(f) \doteq -\log \tilde{Q}(y)$$

$$s.\text{grad}[e, \tilde{y}] = \sum_i -\partial \log P_s[i = \hat{y}[i] \mid \hat{y}] / \partial s[e, \tilde{y}]$$

## Pseudolikelihood Theorem

$$\operatorname{argmin}_Q E_{y \sim \text{Pop}} - \log \tilde{Q}(y) = \text{Pop}$$



## Proof I

We have

$$\min_Q E_{y \sim \text{Pop}} - \log \tilde{Q}(y) \leq E_{y \sim \text{Pop}} - \log \widetilde{\text{Pop}}(y)$$

If we can show

$$\min_Q E_{y \sim \text{Pop}} - \log \tilde{Q}(y) \geq E_{y \sim \text{Pop}} - \log \widetilde{\text{Pop}}(y)$$

Then the minimizer (the argmin) is Pop as desired.

## Proof II

We will prove the case of two nodes.

$$\begin{aligned} & \min_Q E_{y \sim \text{Pop}} - \log Q(y[1]|y[2]) Q(y[2]|y[1]) \\ & \geq \min_{P_1, P_2} E_{y \sim \text{Pop}} - \log P_1(y[1]|y[2]) P_2(y[2]|y[1]) \\ & = \min_{P_1} E_{y \sim \text{Pop}} - \log P_1(y[1]|y[2]) + \min_{P_2} E_{y \sim \text{Pop}} - \log P_2(y[2]|y[1]) \\ & = E_{y \sim \text{Pop}} - \log \text{Pop}(y[1]|y[2]) + E_{y \sim \text{Pop}} - \log \text{Pop}(y[2]|y[1]) \\ & = E_{y \sim \text{Pop}} - \log \widetilde{\text{Pop}}(y|x) \end{aligned}$$

## Contrastive Divergence

**Algorithm (CDk):** Run  $k$  steps of MCMC for  $P_s(\hat{y})$  **starting from**  $y$  to get  $\hat{y}$ .

Then set

$$s.\text{grad}[e, \tilde{y}] = \mathbb{1}[\hat{y}[e] = \tilde{y}] - \mathbb{1}[y[e] = \tilde{y}]$$

**CD Theorem:** If  $P_s(\hat{y}) = \text{Pop}$  then

$$E_{y \sim \text{Pop}} \mathbb{1}[\hat{y}[e] = \tilde{y}] - \mathbb{1}[y[e] = \tilde{y}] = 0$$

**Here we can take  $k = 1$  — no mixing time required.**

**END**